



Docker

Administration

Maruthi S. Inukonda

09th Mar 2019



Agenda

- Containers
- Dockers
- Docker storage
- Docker Identity and Access management
- Docker networking
- Docker Resource (CPU, RAM) allocation
- Container security concerns.



Containers



What is a Container? (1/2) (recap)

- Linux Containers (LXC) is an operating-system-level virtualization method.
- For running multiple isolated Linux systems (containers) on a control host using a single Linux kernel.
- Directly runs on hardware. (No per-instruction level trapping)
- An unprivileged user on host can be privileged user on guest.
- Civil Engineering example :
 - Hostel complex having multiple rooms with shared resources.
Viz., all the above but not study/bedroom.

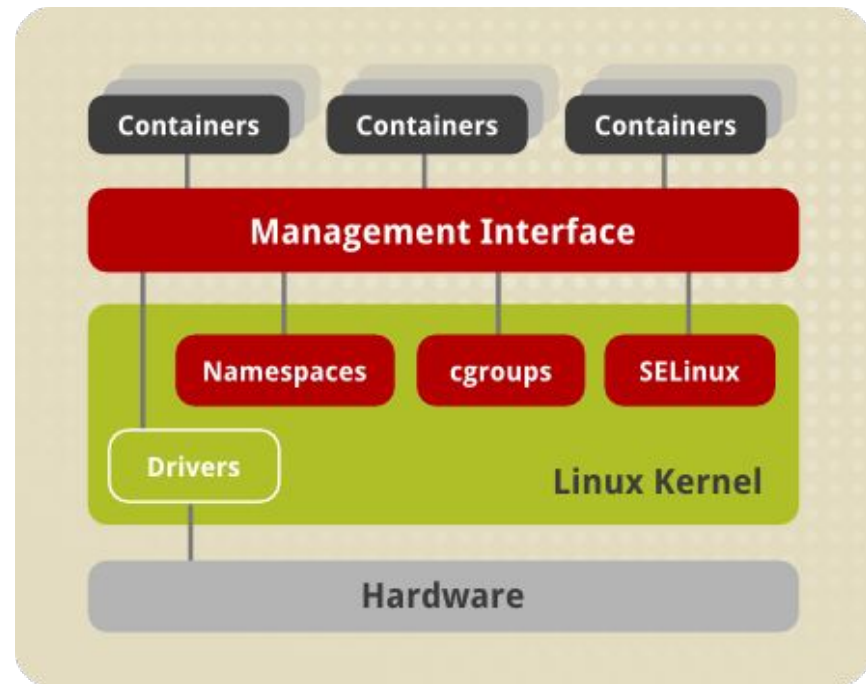


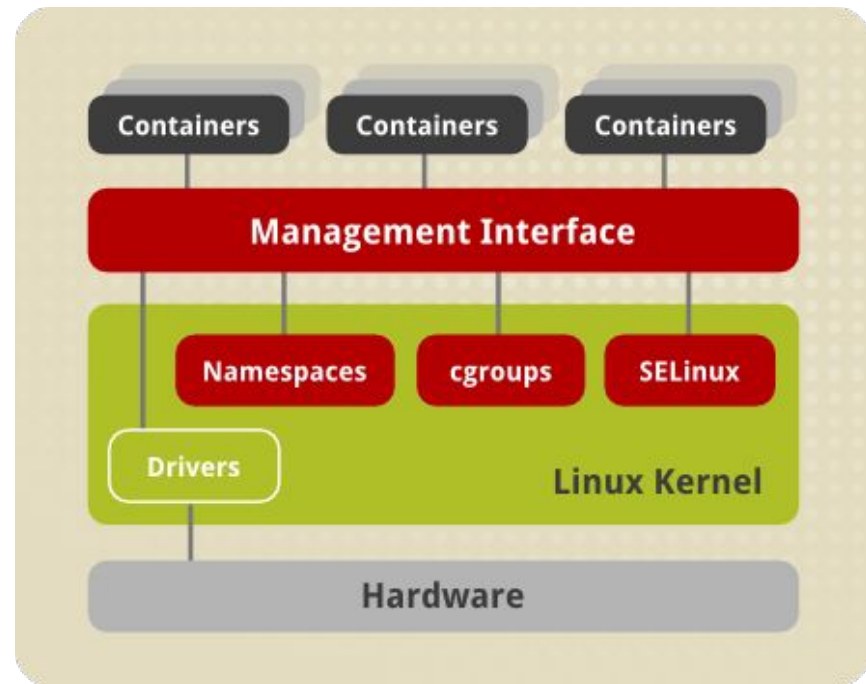
Image Courtesy: Redhat Customer Portal



What is a Container? (2/2) (recap)

It is implemented using following features in Linux

- Advanced Multi-layer Union FS (AUFS) or Overlay FS
- Kernel namespaces
- Cgroups
- Capabilities
- Netfilter, Netlink
- Bind mount
- Role-Based Access Control (RBAC)
 - Eg. SELinux, AppArmor

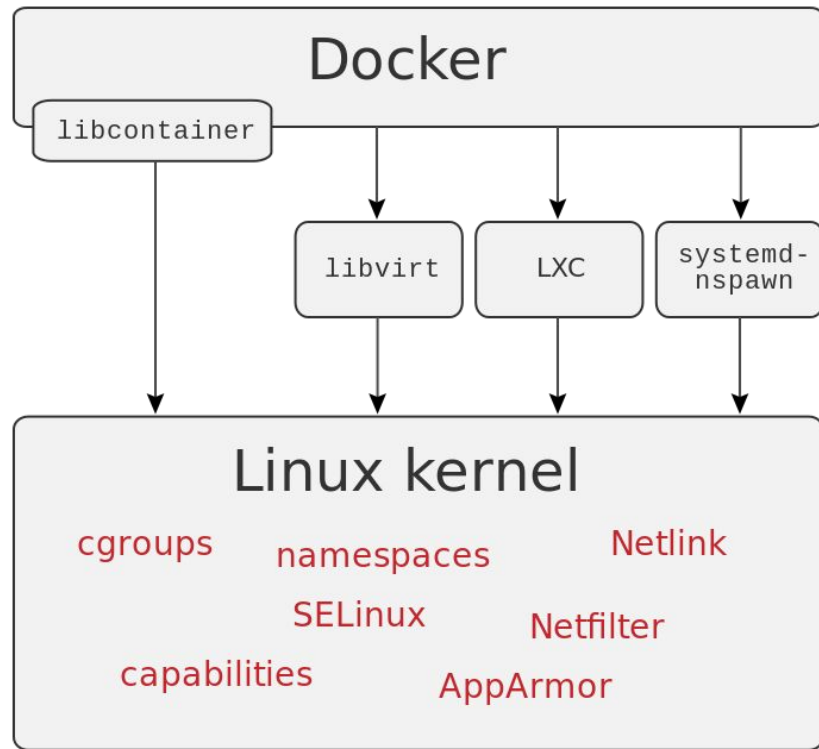


Dockers



What is a Docker?

- Docker is a company that provides software (also called Docker) that allows you to build, run and manage software containers.
- It makes container deployment and administration quite easy.
- It allows re-use of containers created by others.
- It allows running multiple versions of application software with its dependencies on same host.



Docker components

Docker has three major components

- Docker repository/hub
- Docker image
- Docker container



Image Courtesy: Wikipedia



Docker hub

Browser address bar: <https://hub.docker.com/search?isAutomated=0>

Search bar: Search

Notification: 12


Header banner: Docker Store is the new place to discover public Docker content. [Check it out →](#)

Navigation bar:

- Search: nvidia
- Dashboard
- Explore
- Organizations
- Create ▾
- User: maruthi1 ▾

Repositories (774)

Filter: All ▾

	nvidia/cuda public	299 STARS	1M+ PULLS	➤ DETAILS
---	-----------------------	--------------	--------------	-------------------------------



Pulling images

- To pull/download docker images, use
`docker pull <image>:<tag>`

```
$ docker pull ubuntu:latest
latest: Pulling from library/ubuntu
b234f539f7a1: Pull complete
...
d056eaf3dfff4: Pull complete
7dc790b5527b: Pull complete
89bce857a556: Downloading [=====>] 463.4
MB/580.1 MB89bce857a556: Pull complete
Digest: sha256:3cdf1b5becfde8772e15dab594bc76de1cbbefd6c0f8533748854ab47e109ad1
Status: Downloaded newer image for ubuntu:latest
```



Docker container life-cycle

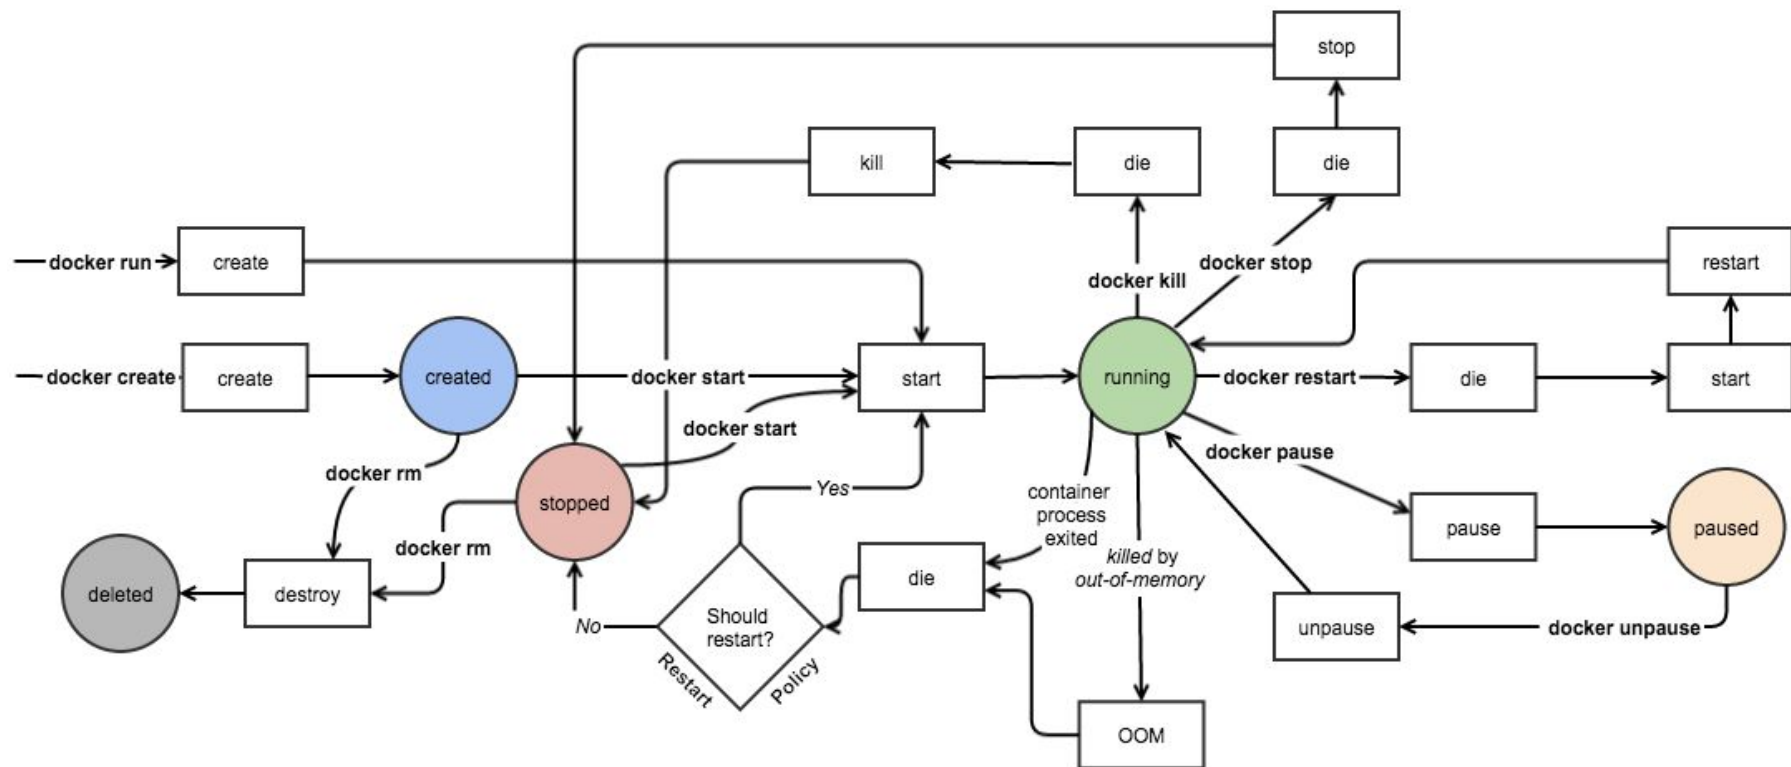


Image Courtesy: Nitin Agarwal @
medium.com

Listing images

- To list downloaded images, use
`docker images`

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu	16.04	a51debf7e1eb	2 weeks ago	116MB
ubuntu	xenial	a51debf7e1eb	2 weeks ago	116MB
ubuntu	trusty	f17b6a61de28	2 weeks ago	188MB
ubuntu	latest	93fd78260bd1	2 weeks ago	86.2MB
nvidia/cuda	9.0-base	74f5aea45cf6	3 weeks ago	134MB
centos	latest	75835a67d134	8 weeks ago	200MB



Running a container

- To run a container, use

```
docker run -it --name <name> <image>:<tag> <program>
```

```
$ docker run -it --name centos1 centos:latest bash  
[root@e7f7395af134 /]#
```

```
[root@e7f7395af134 /]# cat /etc/redhat-release  
CentOS Linux release 7.5.1804 (Core)
```

```
[root@e7f7395af134 /]#
```



Listing running containers

- To list running container, use
`docker ps`

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	8 seconds ago	Up 7 seconds		centos1



Exiting from a container after stopping

- To stop and exit from a container, use `exit`

```
[root@e7f7395af134 /]# exit
```

```
$
```



Listing all containers

- To list all (running/exited) container, use
`docker ps -a`

```
$ docker ps -a
CONTAINER ID   IMAGE          COMMAND         CREATED        STATUS        PORTS          NAMES
e7f7395af134   centos:latest  "bash"         8 seconds ago  Exited (0)    5 seconds ago
centos1
```



Exiting from a container without stopping

- To exit from a container without stopping, press
`Ctrl+p Ctrl+q`

```
[root@e7f7395af134 /]# Ctrl+p Ctrl+q      read escape sequence
```

```
$
```



Creating a container

- To create a container, use

```
docker create --name <name> -it <image>:<tag> <program>
```

```
$ docker create --name centos2 -it centos:latest bash
916dc303760db834c1aa4a9b591605ab56c91aba97bceb6fda2ca0db564f489a
```

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	31 seconds ago	Created		centos1
916dc303760d	centos:latest	"bash"	About a minute ago	Created		centos2



Starting a container

- To start a container, use
`docker start <name_or_id>`

```
$ docker start centos2
centos2
```

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e7f7395af134	centos:latest	"bash"	4 minutes ago	Up 3 minutes		centos1
916dc303760d	centos:latest	"bash"	3 minutes ago	Up 2 minutes		centos2



Attaching to a running container

- To start a container, use
`docker attach <name_or_id>`

```
$ docker attach centos2  
[root@916dc303760d /]#
```



Storage



Sharing a host directory/file to container

- To share a host directory or file to a container, use

```
docker run -v <host_dir>:<guest_dir> <image>:<tag>  
<program>
```

```
$ docker run -v /mnt:/mnt -it ubuntu:xenial bash
```



Identity and Access



Sharing host credentials with container

- To share a host credentials with a container, use

```
docker run -u `id -u`:`id -g` -v \  
    /etc/passwd:/etc/passwd <image>:<tag> <program>
```

```
$ docker run -ti -u `id -u`:`id -g` --name "ubuntu-dock" -v /etc/passwd:/etc/passwd -v  
/home/maruthisi:/home/maruthisi ubuntu:xenial bash
```



Networking



Listing software defined networks (SDN)

- To list SDNs

```
docker network ls
```

```
$ ifconfig
```

```
docker0  Link encap:Ethernet  HWaddr 02:42:f7:e7:61:9e  
         inet addr:172.17.0.1   Bcast:172.17.255.255   Mask:255.255.0.0
```

```
...
```

```
enp0s25  Link encap:Ethernet  HWaddr 90:1b:0e:e5:90:3e  
         inet addr:192.168.136.108  Bcast:192.168.136.255   Mask:255.255.255.0
```

```
...
```

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
70176b89a07a	bridge	bridge	local
a8e63013eb97	host	host	local



Using host network

- To launch an ubuntu docker with host networking, use
`docker run -it --net host <image>:<tag> <command>`

```
$ docker run -it --net host ubuntu:xenial bash
```



Using bridge network

- To launch an ubuntu docker with bridge networking, use
`docker run -it --net bridge <image>:<tag> <command>`

```
$ docker run -it --net bridge ubuntu:xenial bash
```



Creating custom bridge network

- To create a custom bridge network and router

```
docker network create --subnet=<network>/<prefix> <name>
```

```
$ docker network create --subnet=172.18.0.0/16 mysdn1
```

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
70176b89a07a	bridge	bridge	local
a8e63013eb97	host	host	local
f85eb2d101ef	mysdn1	bridge	local
b9550c4a7b89	none	null	local



Using custom bridge network

- To launch an ubuntu docker with host networking, use
`docker run -it --net <name> <image>:<tag> <command>`

```
$ docker run -it --net mysdn1 ubuntu:xenial bash
# ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:ac:12:00:02
          inet addr:172.18.0.2  Bcast:172.18.255.255  Mask:255.255.0.0
          ...
```



“apt update” and “apt install net-tools” inside the container to install ifconfig



Resource Allocation



Limiting RAM usage

- Using cgroups docker containers could be reserved/restricted resources.
- Use `--memory` to limit the memory (RAM) usage.
- Use `--memory-swap` to limit the virtual memory (RAM+Swap) usage.

```
$ docker run -it --memory=512m --memory-swap=1g ubuntu:latest bash
```



Limiting CPU usage

- Use `--cpuset-cpus` to restrict the usable CPUs.

```
$ docker run -it --cpuset-cpus=0,2 ubuntu:latest bash
```



Container security concerns



Privilege Escalation

- A container started by an unprivileged user can get root shell on the host.
 - Every container is started as a child process of the daemon running as root.
 - Child processes also run as root

On the host

```
$ cat /etc/hostname
```

Launch a guest and modify the file

```
$ docker run -it -v /etc/hostname:/etc/hostname --name ub1 ubuntu:latest bash
# id
# vi /etc/hostname
```

On the host

```
$ cat /etc/hostname
```



Lack of isolation

- A container started by an unprivileged user can be attached by another user. This is not acceptable in multi-tenant environment.

From one user

Launch a guest and modify the file

```
$ id
$ docker run -it --name ub2 ubuntu:latest bash
# top
```

From another user

```
$ id
$ docker attach ub2
```



Solutions/Workarounds

Solutions to the isolation and privilege escalation.

- Docker Enterprise Edition uses Role Based Access Control (RBAC)
 - Container daemon runs as root. A child process (called proxy daemon) is created with the same RBAC context as the user.
 - Every container is started as a child process of the proxy-daemon.
- AWS uses docker community edition, but runs containers inside VMs to provide isolation.

Workaround to the privilege escalation

- Idmapping technique is used to map host's uid/gid to non-existent uid/gid in guest.



Q & A

